



NETWORK STRUCTURE & CLOUD COMPUTING PROJECT REPORT

“SPPARD ON CLOUD”

APRIL 2016

*Guided By-
Prof. Julian Cooper*

*By-
Aayush Shah
Dishu Jindal
Palash Kochar
Pawan Matta
Rimple Talati
Saurabh Goyal*

Contents

| | |
|---|----|
| INTRODUCTION..... | 3 |
| Types of Cloud | 3 |
| REQUIREMENTS..... | 4 |
| BUSINESS JUSTIIFICATION | 4 |
| SYSTEMS AND DESCRIPTIONS..... | 5 |
| Platform..... | 5 |
| Techniques Used | 5 |
| Web Application | 5 |
| CLOUD ARCHITECTURE..... | 8 |
| Visio Diagram..... | 8 |
| Virtual Private Cloud..... | 9 |
| Amazon RDS for MySQL..... | 9 |
| Elastic Load Balancer | 10 |
| Web Application Security | 11 |
| Auto Scaling | 11 |
| Cloud Watch..... | 12 |
| Launching an EC2 instance..... | 13 |
| Simple Storage Service (S3)..... | 14 |
| Cloud Pricing Model for the next 3-6 months | 15 |
| Disaster Recovery Procedures | 16 |
| References..... | 17 |

INTRODUCTION

Cloud computing is Internet based development and use of computer technology. It is a paradigm shift where details are abstracted from the users who no longer need knowledge or expertise or control over the technology infrastructure "in the cloud" that supports them. It typically involves the provision of dynamically scalable and often virtualized resources as a service over the Internet.

'Cloud' refers to a distinct IT environment that is designed for the purpose of remotely provisioning scalable and measured IT resources. The term cloud is originated as a metaphor for the Internet. It is important to note the differences between the term 'cloud' and cloud symbol from Internet. There are multiple individual clouds that has a finite boundary and are accessible via the Internet. On one hand, where Internet provides open access to many web based IT resources, a cloud, on other hand typically privately owned offers access to IT resources that is metered.

Cloud computing typically provides 3 types of services: Infrastructure as a Service (IaaS), Platform as a Service (PaaS) and Software as a Service (SaaS). These services are available over the Internet to the whole world where the cloud acts as a single point of access for serving all the customers.

Types of Cloud

- Private Cloud – This is solely created for the internal purposes for a single organization. It is typically the first step of the organization to move into the cloud.
- Public Cloud – Multiple organizations rents the clouds services from cloud provider on-demand basis.
- Hybrid Cloud – This is composed of multiple internal or external cloud. For example, a cloud consumer may choose to deploy sensitive data to a private and less sensitive cloud services to a public cloud.
- Community Cloud – It is similar to public cloud except that its access is limited to a specific community of cloud consumers.

REQUIREMENTS

- Entire stack has to be on one either AWS or Azure.
- Java Spring developed login portal secured by self signed certificate
- Backend DB has to be either relational DB or NoSQL with sample data
- Java spring developed use case buttons
 - First button - creation of 3 user login to application and run 2 unique report
 - Second button - creation of 10 user login to application and run 6 unique report
 - Third button - creation of 17 user login to application and run 10 unique report
 - Fourth button - creation of 24 user login to application and run 14 unique report
 - Fifth button - user input fields to allow the ability to enter parameters for how many users login creation and how many reports those users are uniquely are running.
- Proper infrastructure alerts and triggers to allow for auto-scaling of resources to accommodate the additional load in application, network, data storage and usage with your environment

BUSINESS JUSTIIFICATION

To deploy a Spring MVC web application in AWS with the help of MySQL RDS database to scale in/out the server according to user traffic and load and hence avoiding latency and bottlenecks.

SYSTEMS AND DESCRIPTIONS

Platform

- Amazon Web Services

Typically when we need to choose a cloud based platform then only two of them comes in the mind – Microsoft’s Azure or Amazon’s AWS.

As a student, we had free credits for both of these platforms. So the question arises, which platform will be best suitable for our project?

When we started to research on these two platforms, we observed following comparisons:

| Criteria | AWS | Azure |
|----------------------|----------------------------------|-------------------------------------|
| Instance Families | 7 | 4 |
| Instance Types | 38 | 33 |
| Zones | Yes | - |
| Ephemeral(Temporary) | Yes | Temporary storage – D Drive |
| Block Storage | EBS | Page Blobs |
| Relational DB | RDS | Relational DBs |
| NoSQL and Big Data | DynamoDB, EMR, Kinesis, Redshift | Windows Azure table, HDInsight |
| Virtual Network | VPC | VNet |
| DNS | Route53 | - |
| Pricing | Per hour – rounded trip | Per minute – rounded up commitments |

We had used both of them during our assignments and apart from the above comparisons we also felt that the performance speed of Azure is little slow as compared to AWS. So we decided to build our entire project on AWS.

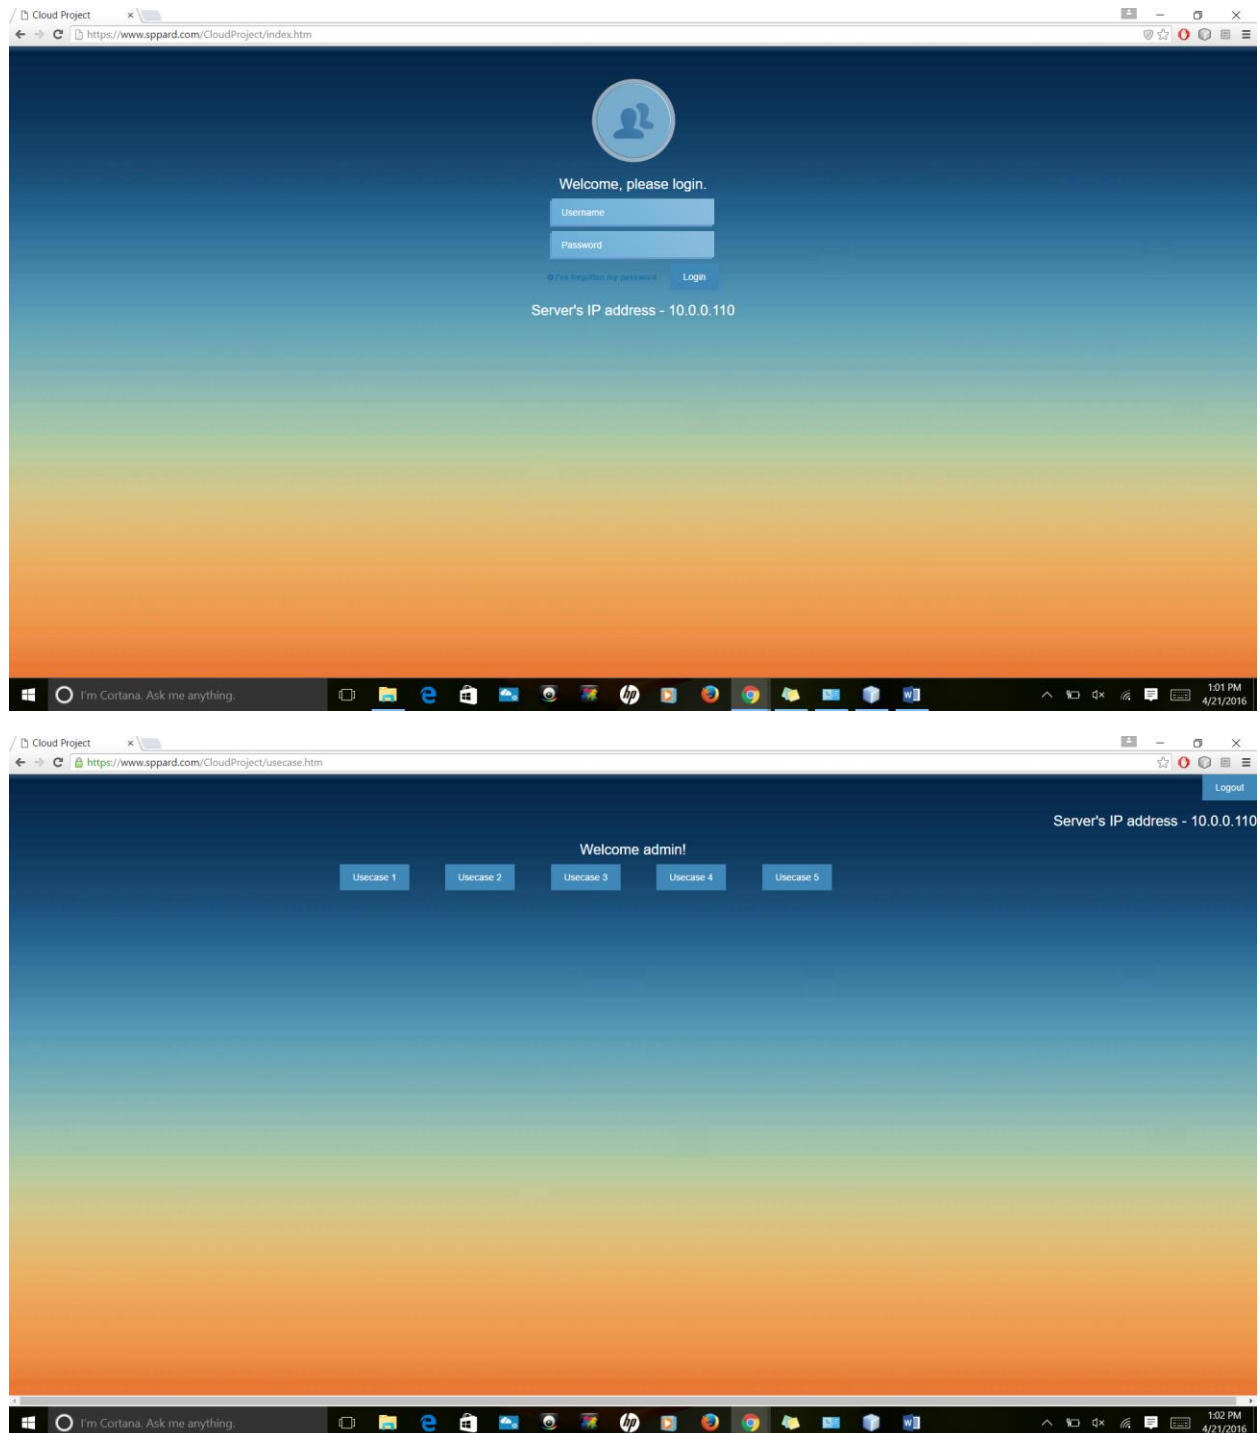
Techniques Used

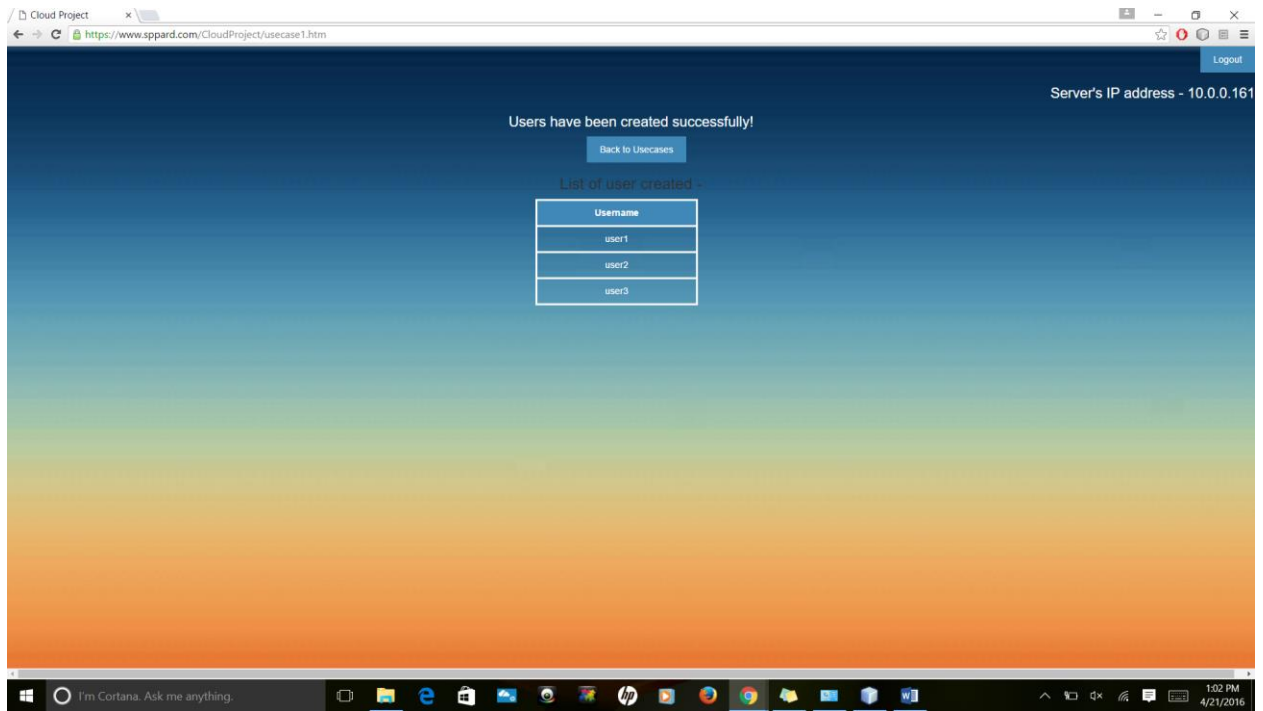
- Elastic Load Balancing
- Auto Scaling
- Cloud Watch
- Health Check

Web Application

- We are creating an E-commerce application <https://www.sppard.com/> based on Spring MVC framework
- We are using MySQL with RDS as a database
- We created one table which consists of userID, username, passwords, product_names. The random data was filled by - <https://www.devart.com/dbforge/sql/data-generator/>

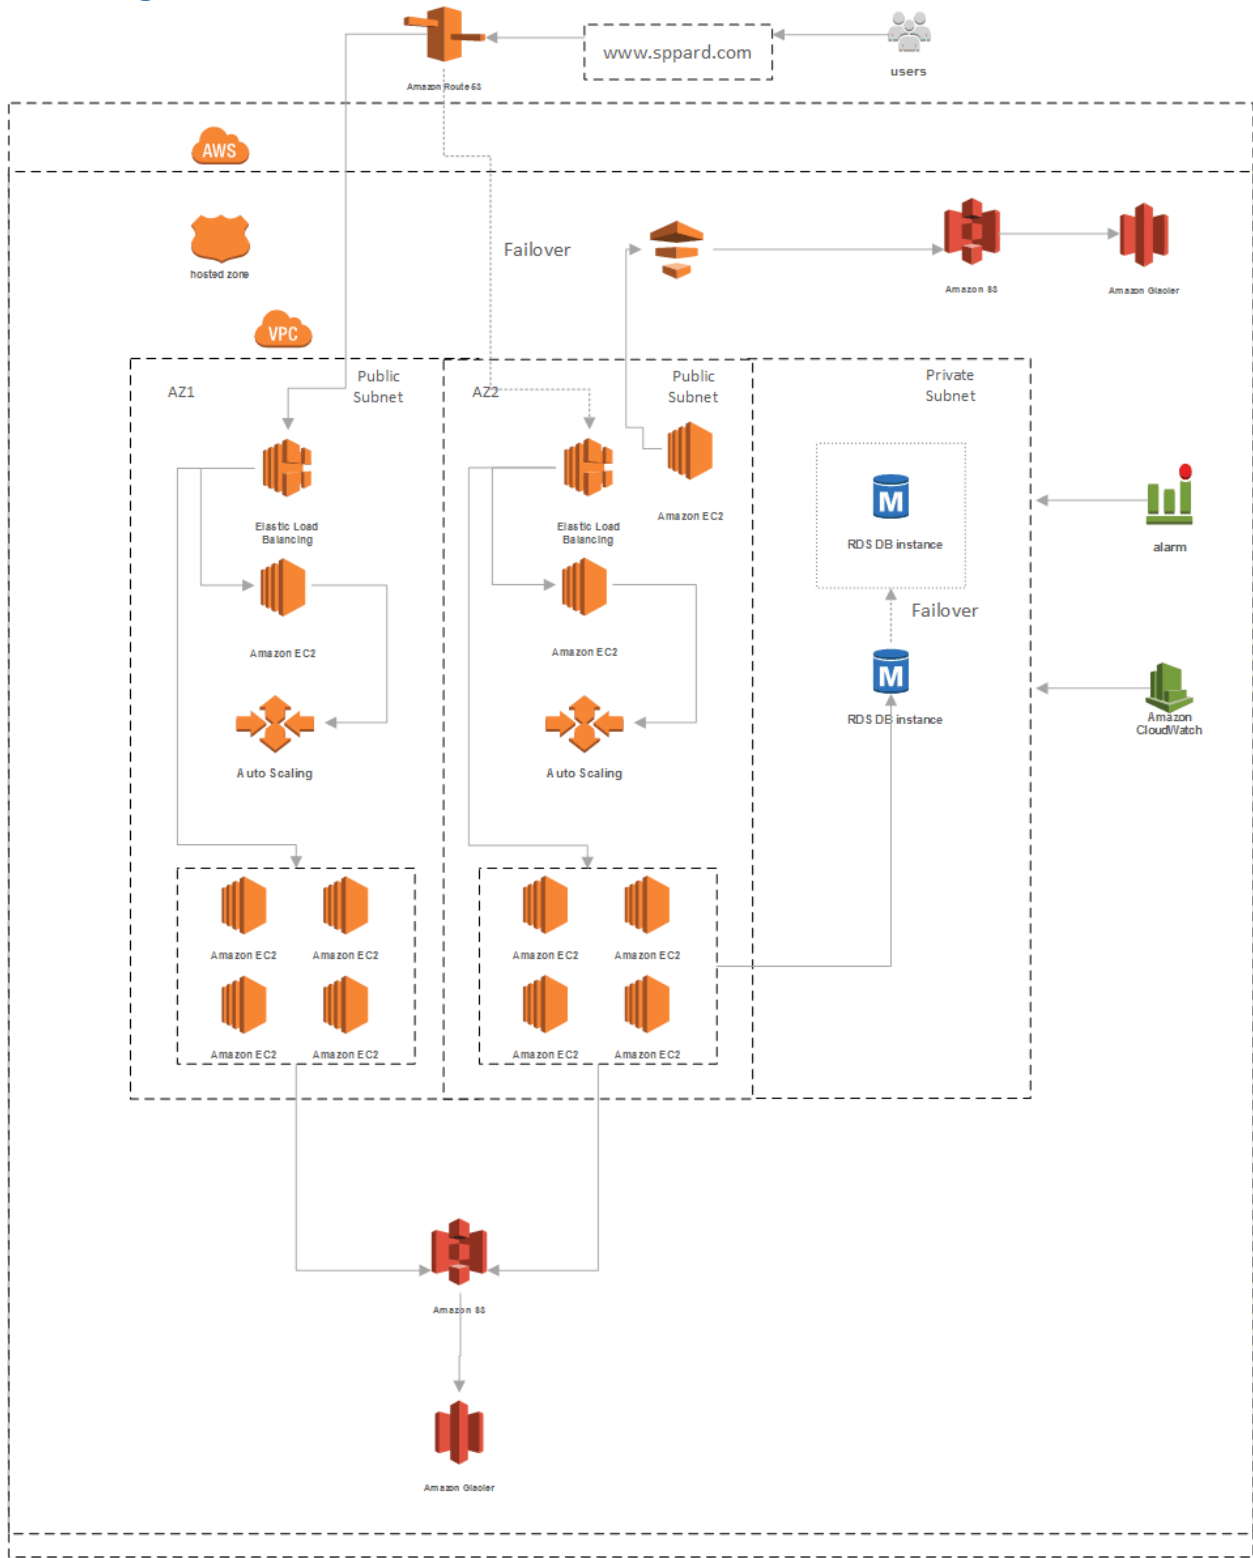
- Size of the dataset is 760MB and it has around 16 million records.
- We have written a Java code which allows to create certain number of users on click of Use Case buttons. We are also storing these created users in the session which shows up in the next page. Besides user creation, we are also generating certain number of unique reports on each Use case button and these reports are automatically saved on S3 bucket in AWS.





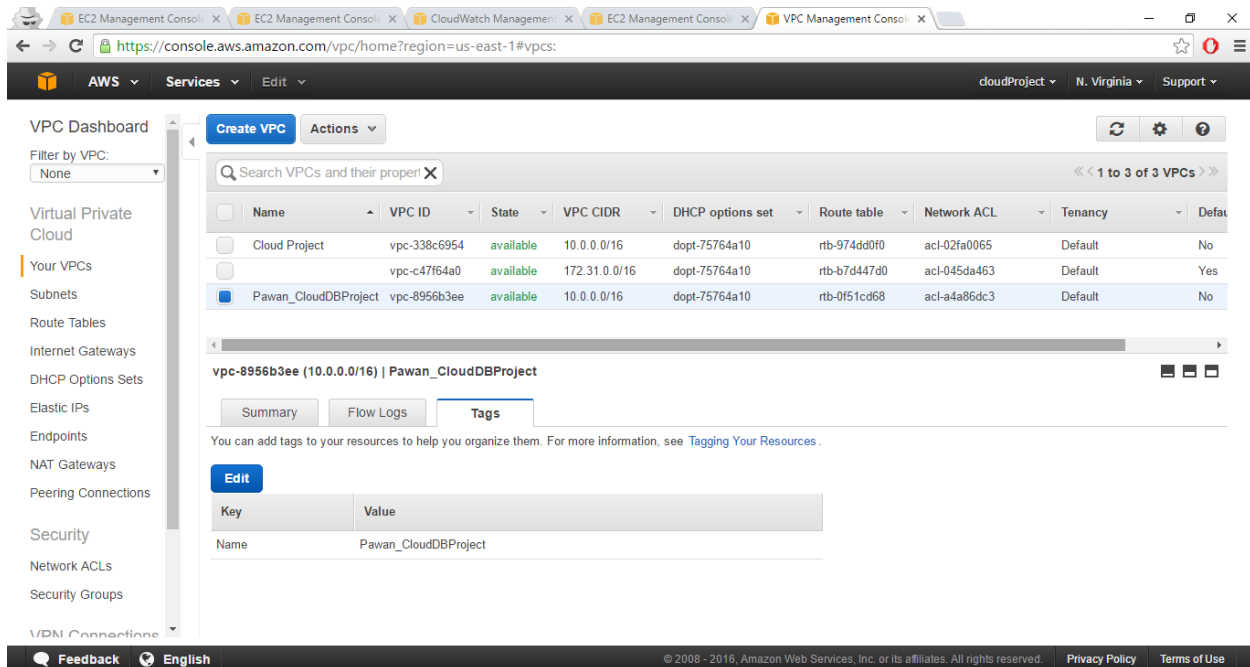
CLOUD ARCHITECTURE

Visio Diagram



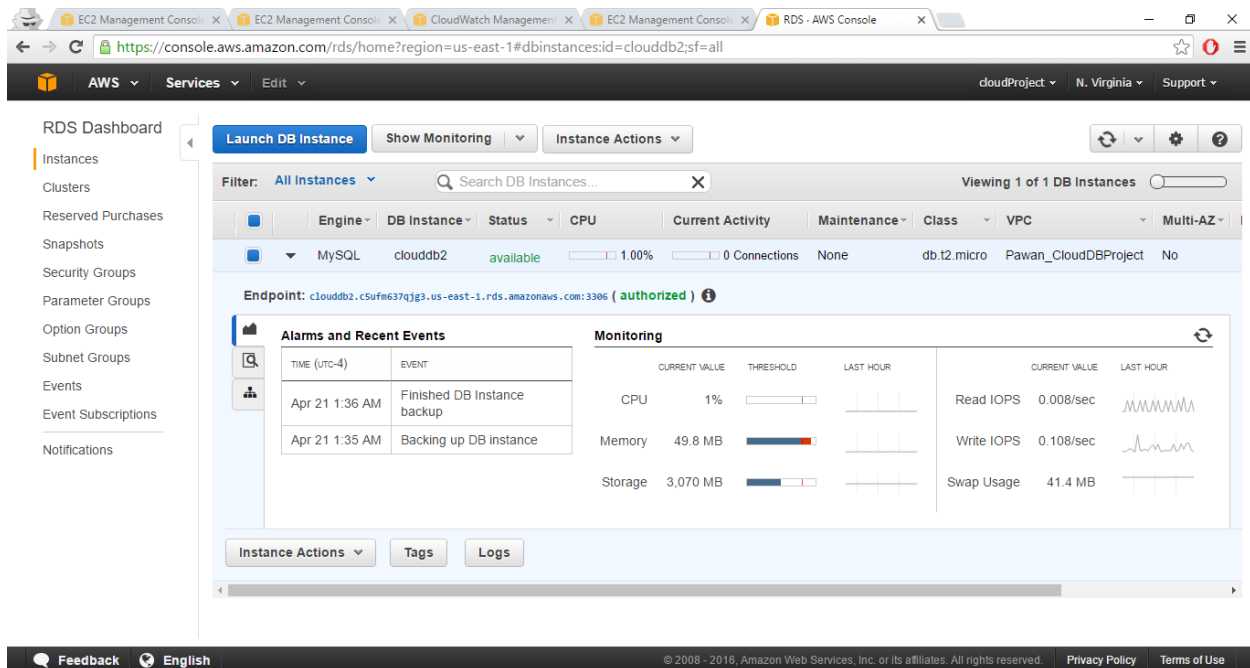
Virtual Private Cloud

- Virtual private cloud (VPC) is a virtual network dedicated to our AWS account. It is logically isolated from other virtual networks in the AWS cloud.
- We had created VPC with Public and Private Subnets.
- EC2 instances, RDS instance, Load Balancer, Auto Scaling Group and security groups are inside VPC



Amazon RDS for MySQL

- Amazon RDS makes it easy to set up, operate, and scale MySQL deployments in the cloud.
- It has following features –
 - Scalable
 - Cost-efficient
 - Resizable hardware capacity
- We had launched the RDS instance in the private subnet of VPC in order to secure the data.
- During configuration we had selected MYSQL, specified DB details and db.t2.micro as DB instance class. On click of configured advanced settings, we selected the VPC that we have created and gave other information.



RDS connection with MySQL workbench

- As the RDS is in private subnet so no one access it from outside. In order to load the data in to RDS, we have created a VPN instance in our VPC through which we can connect the RDS instance.
- In order to create the VPN instance, we had used an AMI from Amazon marketplace and used openVPN to connect to the VPN instance.

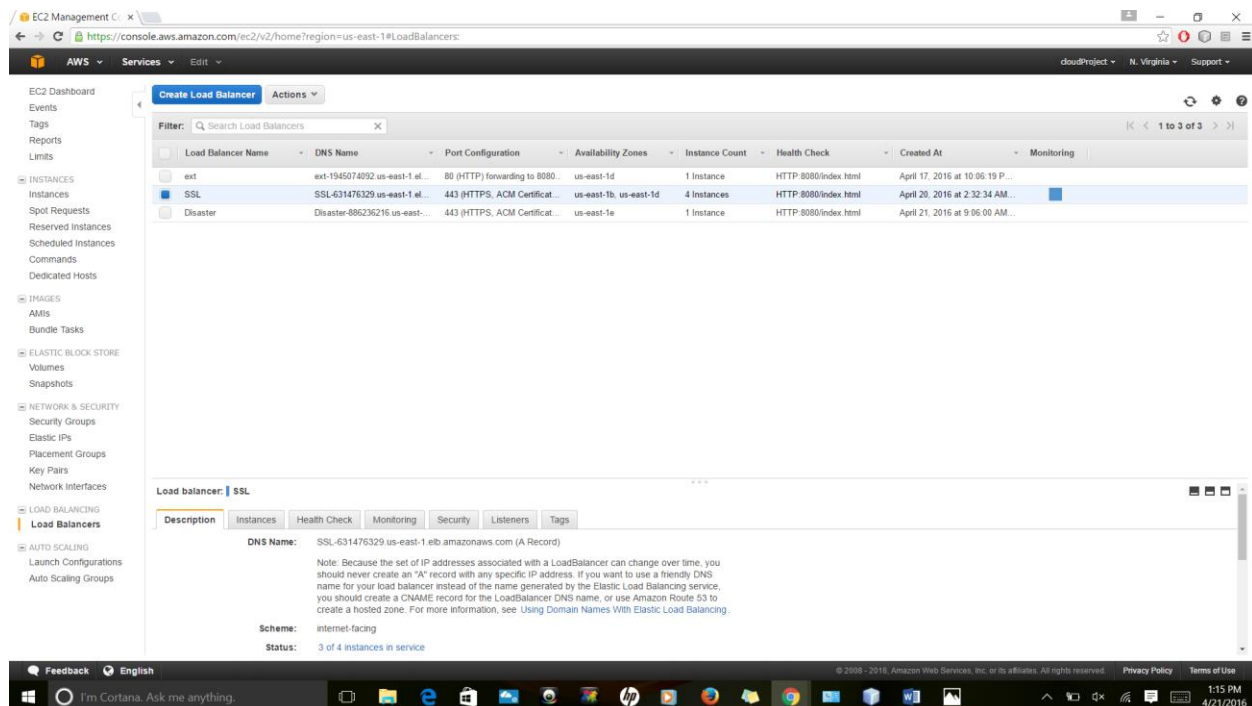
Advantages of VPN –

- VPN offers secure connection to a remote network via the Internet.
- These are very useful for connecting multiple networks together securely.
- VPN encrypts Internet traffic which helps to stymie other people who may be trying to capture your passwords

Elastic Load Balancer

- It is a load balancing solution that automatically scales its request-handling capacity in response to incoming application traffic.
- ELB solutions often provide scalable cloud computing capacity.
- It has following features –
 - Detection of unhealthy instances
 - Spreading instances across healthy channels only.
 - Flexible cipher support.
 - Centralized management of SSL certificates.
- As our application is secured through self-signed certificate, we have to attach SSL certificate in load balancer so that it can forward HTTPS request to our EC2 instances hosting the web application. We need to add listener by:-

- LB Protocol – HTTPS
- LB Port – 443
- Instance Protocol – HTTPS



Web Application Security

- To secure our application, we have used self-signed certificate.
- We had uncommented the connector port – 8443 and added keystore file and keystore pass so that it accepts HTTPS requests.
- We had done changes to redirect 8080 to 8443 in server.xml and added security constraint in application's web.xml file.

Auto Scaling

- Auto Scaling allows us to scale our Amazon EC2 capacity up or down automatically according to conditions we define.
- In the launch configuration, we have used pre-configured AMI of EC2 instance that had tomcat server and war file of our application.
- We are using m3.large instance class for the EC2 instances that will be created.
- Next, we added security group that will control the traffic of our EC2 instances.
- Now, we had created auto scaling group. We set the minimum as 1 and maximum as 3 instances.
- We have used Cloud Watch to monitor our server instances and at certain event like CPU utilization above or below the specified limit, it will automatically launch or decrease instances based on the AMI template that we define in launch configurations. These instances are attached to the Load Balancer which have been added in this ASG.

The first screenshot shows the 'Details' tab for the Auto Scaling Group 'autoScale-3'. The table below summarizes the configuration:

| Property | Value |
|---------------------------|---------------|
| Name | autoScale-3 |
| Launch Configuration | micro-working |
| Instances | 2 |
| Desired | 1 |
| Min | 1 |
| Max | 3 |
| Availability Zones | us-east-1d |
| Default Cooldown | 300 |
| Health Check Grace Period | 298 |

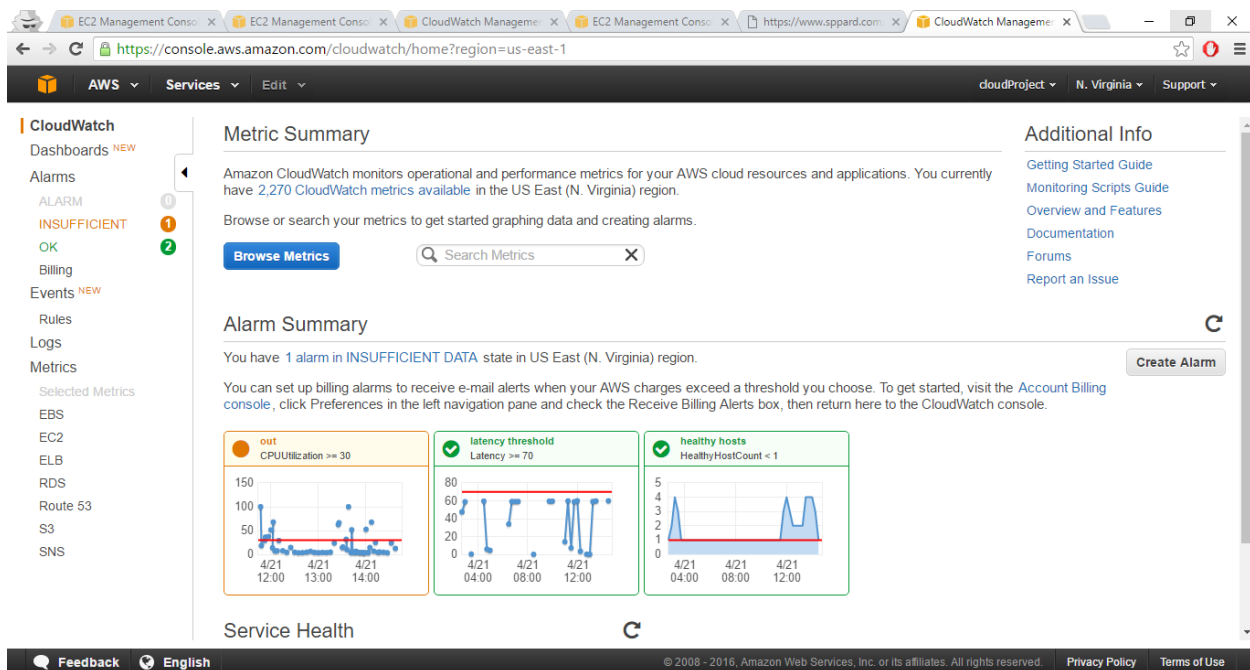
The second screenshot shows the 'Activity History' tab for the same Auto Scaling Group. The table below lists the recent activities:

| Status | Description | Start Time | End Time |
|-------------|---|------------------------------|------------------------------|
| In progress | Terminating EC2 Instance: i-00b5b2d0c1a69064f | 2016 April 21 13:22:04 UTC-4 | |
| Successful | Launching a new EC2 Instance: i-04429c2d6106d1c | 2016 April 21 13:17:56 UTC-4 | 2016 April 21 13:19:00 UTC-4 |
| Successful | Terminating EC2 Instance: i-093a2ecbcb2a2fe | 2016 April 21 13:15:51 UTC-4 | 2016 April 21 13:17:25 UTC-4 |
| Successful | Launching a new EC2 Instance: i-0c409176722d40c2 | 2016 April 21 13:12:18 UTC-4 | 2016 April 21 13:23:23 UTC-4 |
| Successful | Terminating EC2 Instance: i-0fc3ac9207c3b3a0 | 2016 April 21 13:11:11 UTC-4 | 2016 April 21 13:12:42 UTC-4 |
| Successful | Launching a new EC2 Instance: i-00b5b2d0c1a69064f | 2016 April 21 13:06:12 UTC-4 | 2016 April 21 13:07:54 UTC-4 |
| Successful | Terminating EC2 Instance: i-091f11ec66e795ee | 2016 April 21 13:03:55 UTC-4 | 2016 April 21 13:05:52 UTC-4 |
| Successful | Launching a new EC2 Instance: i-093a2ecbcb2a2fe | 2016 April 21 13:01:53 UTC-4 | 2016 April 21 13:03:18 UTC-4 |
| Successful | Launching a new EC2 Instance: i-0fc3ac9207c3b3a0 | 2016 April 21 12:59:20 UTC-4 | 2016 April 21 13:00:28 UTC-4 |
| Successful | Terminating EC2 Instance: i-0258b66a1741b7e3 | 2016 April 21 12:57:10 UTC-4 | 2016 April 21 12:58:59 UTC-4 |
| Successful | Launching a new EC2 Instance: i-091f11ec66e795ee | 2016 April 21 12:52:01 UTC-4 | 2016 April 21 12:53:07 UTC-4 |
| Successful | Terminating EC2 Instance: i-0b4d54951c332a35e | 2016 April 21 12:48:53 UTC-4 | 2016 April 21 12:50:48 UTC-4 |
| Successful | Terminating EC2 Instance: i-02e7a8f9ed8a21202 | 2016 April 21 12:46:47 UTC-4 | 2016 April 21 12:47:16 UTC-4 |
| Successful | Launching a new EC2 Instance: i-0258b66a1741b7e3 | 2016 April 21 12:44:28 UTC-4 | 2016 April 21 12:45:35 UTC-4 |
| Successful | Terminating EC2 Instance: i-0f35b18d34b39139 | 2016 April 21 12:40:05 UTC-4 | 2016 April 21 12:41:27 UTC-4 |
| Successful | Terminating EC2 Instance: i-0f35b18d34b39139 | 2016 April 21 12:40:05 UTC-4 | 2016 April 21 12:41:27 UTC-4 |

Cloud Watch

- Amazon CloudWatch is a component of Amazon Web Services (AWS) that provides monitoring for AWS resources such as Amazon EC2 instances, Load Balancers and Amazon RDS instances and the customer applications running on the Amazon infrastructure.

- We had set the Alarm to monitor the CPU usage of our EC2 instances and then used this data to determine whether to launch additional instances to handle increased load
- An alarm can have three possible states:
 - OK—The metric is within the defined threshold
 - ALARM—The metric is outside of the defined threshold
 - INSUFFICIENT_DATA—The alarm has just started, the metric is not available, or not enough data is available for the metric to determine the alarm state
- We had set 1 alarm to trigger:
 - If the CPU utilization is greater than 30% for 1 minute then ASG will scale out by 1 instance.



Launching an EC2 instance

- First we launched an EC2 instance, installed TOMCAT and deployed the war file of our application and created its AMI. This AMI will be used by ASG to launch new instances.

The screenshot shows the AWS Management Console for EC2 Instances. The left sidebar contains navigation links for EC2 Dashboard, Events, Tags, Reports, Limits, INSTANCES, INSTANCES (selected), Spot Requests, Reserved Instances, Scheduled Instances, Commands, Dedicated Hosts, IMAGES, AMIs, Bundle Tasks, ELASTIC BLOCK STORE, Volumes, Snapshots, NETWORK & SECURITY, Security Groups, Elastic IPs, and Placement Groups. The main content area displays a table of instances. The selected instance is 'secondary ASG' with ID 'i-0f52fcc9be7f35329', type 'm3.medium', and state 'running'. The console also shows the 'Status Checks' tab for this instance, indicating that both System and Instance status checks have passed.

| aws:autoscal | Instance ID | Instance Type | Availability Zone | Instance State | Status Checks | Alarm Status | Public DNS | Public IP |
|---------------|---------------------|---------------|-------------------|----------------|----------------|--------------|--------------------------|---------------|
| secondary ASG | i-0f52fcc9be7f35329 | m3.medium | us-east-1e | running | 2/2 checks ... | None | ec2-54-87-210-92.comp... | 54.87.21... |
| scalingGroup | i-0d1c651a8e4a103cf | t1.micro | us-east-1b | terminated | | None | | |
| scalingGroup | i-02f20051432460d10 | t1.micro | us-east-1b | running | 2/2 checks ... | None | ec2-54-152-251-38.com... | 54.152.251... |
| scalingGroup | i-05a29db7543a4bc56 | t1.micro | us-east-1d | terminated | | None | | |

Instance: **i-0f52fcc9be7f35329** Public DNS: **ec2-54-87-210-92.compute-1.amazonaws.com**

Status Checks

Status checks detect problems that may impair this instance from running your applications. [Learn more](#) about status checks.

Create Status Check Alarm

System Status Checks ⓘ

These checks monitor the AWS systems required to use this instance and ensure they are functioning properly.
System reachability check passed

Instance Status Checks ⓘ

These checks monitor your software and network configuration for this instance.
Instance reachability check passed

Simple Storage Service (S3)

- It is an online file storage web service.
- We have used S3 to directly save the generated reports on click of each Use case button.

The screenshot shows the AWS Management Console for S3 Buckets. The left sidebar contains navigation links for S3 Dashboard, Buckets, Objects, and Permissions. The main content area displays a table of buckets. The selected bucket is 'reports-sppard'. The console also shows the 'Properties' tab for this bucket, indicating that the bucket is in the 'US Standard' region and was created on 'Wed Apr 20 05:20:18 GMT-0400 2016'.

| Name | Size |
|----------------|------|
| dbfullbackup | |
| incrementalbu | |
| jsonn | |
| newbucket12343 | |
| reports-sppard | |

Bucket: reports-sppard

Bucket: reports-sppard
Region: US Standard
Creation Date: Wed Apr 20 05:20:18 GMT-0400 2016
Owner: dxtlg

Permissions
 Static Website Hosting
 Logging
 Events
 Versioning
 Lifecycle
 Cross-Region Replication
 Tags
 Requester Pays
 Transfer Acceleration

Cloud Pricing Model for the next 3-6 months

It turns out that estimate cost of 'sppard' application is \$59.98 per month which brings in \$359.88 for 6 months.

Link to Billing Calculator:

<http://calculator.s3.amazonaws.com/index.html#r=IAD&key=calc-1BD522FD-478C-4414-B1FD-1D96EBC45A28>

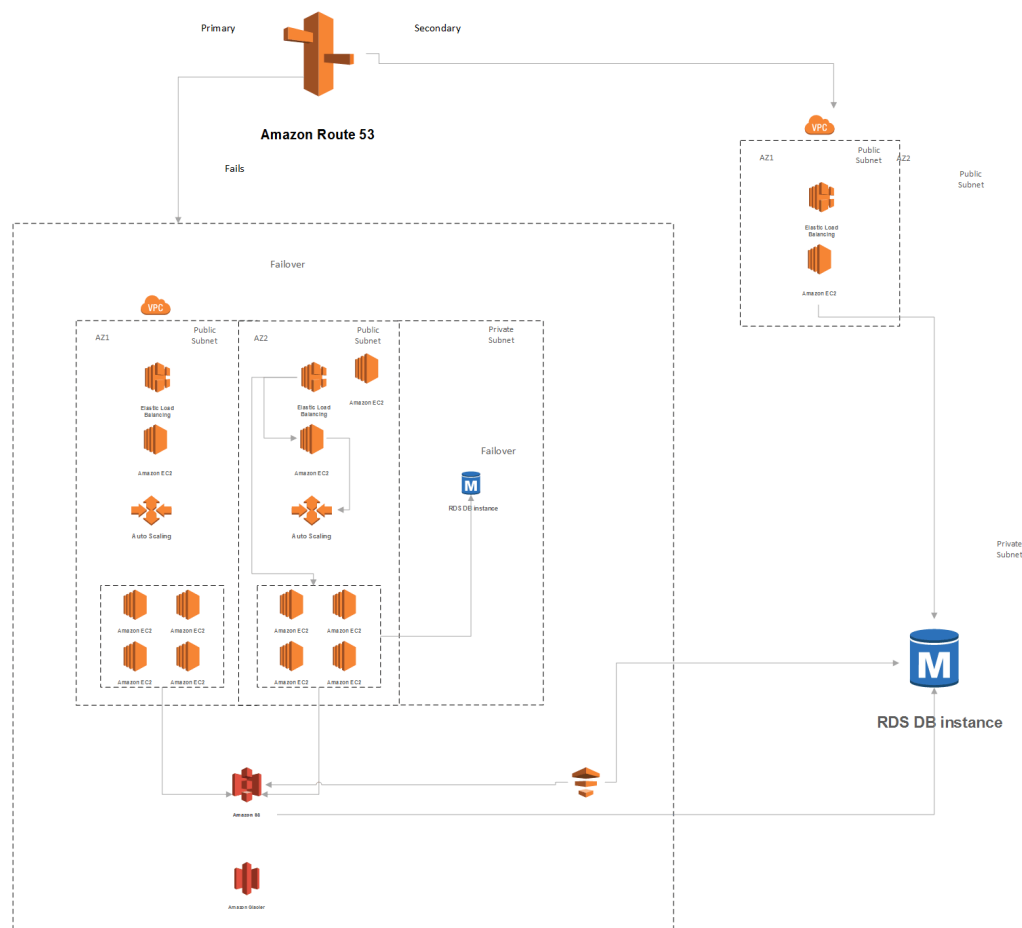
The screenshot displays the AWS Simple Monthly Calculator interface. At the top, it shows the AWS logo and the text 'SIMPLE MONTHLY CALCULATOR'. Below this, there's a navigation bar with 'Services' and 'Estimate of your Monthly Bill (\$ 105.58)'. The main content area is titled 'Estimate of Your Monthly Bill' and includes a 'Show First Month's Bill' checkbox. Below this, a table lists various AWS services and their estimated costs. The services listed are: Amazon EC2 Service (US-East) at \$64.82, Amazon S3 Service (US-East) at \$2.00, Amazon Route 53 Service at \$1.54, Amazon RDS Service (US-East) at \$14.53, Amazon DynamoDB Service (US-East) at \$0.00, Amazon CloudWatch Service (US-East) at \$36.60, Amazon VPC Service (US-East) at \$0.00, AWS Data Transfer In at \$28.71, AWS Data Transfer Out at \$0.00, and AWS Support (Basic) at \$-42.62. The total monthly payment is \$105.58. On the right side, there's a 'Common Customer Samples' section with links to various AWS services and samples.

| Service | Estimated Cost |
|-------------------------------------|-----------------|
| Amazon EC2 Service (US-East) | \$64.82 |
| Amazon S3 Service (US-East) | \$2.00 |
| Amazon Route 53 Service | \$1.54 |
| Amazon RDS Service (US-East) | \$14.53 |
| Amazon DynamoDB Service (US-East) | \$0.00 |
| Amazon CloudWatch Service (US-East) | \$36.60 |
| Amazon VPC Service (US-East) | \$0.00 |
| AWS Data Transfer In | \$28.71 |
| AWS Data Transfer Out | \$0.00 |
| AWS Support (Basic) | \$-42.62 |
| Total Monthly Payment: | \$105.58 |

Disaster Recovery Procedures

Disaster Recovery is a process in which if a primary network or a dedicate server in a zone or a zone in a whole collapses or crashes, there should be a procedure in which we try to bring the whole consistent system as a whole.

Until the system is down, we point the Route53 on another system which is nowhere connected to the earlier primary, system. The entire load which was directed to the primary network or system gets redirected to the new system/zone or the database



Steps for Disaster Recovery Plan

1. First we would be taking the entire backup on a weekly basis and saving that into an S3 bucket using an data pipeline
2. We would then take a incremental backup and then again store that into the S3 bucket with a different name
3. We would create a New SQL which is consistent using this CSV file or utilize the RDS instance which is already available since it's a multi -az

References

<https://aws.amazon.com/documentation/>

<http://searchaws.techtarget.com/>

<http://whatisccloud.com/>